

# Missguided attempts at performance optimizations

Radu Ciorba {radu@devrandom.ro}

February 17, 2017

# What is Borg Backup?

- Deduplicating backup program
- written mostly in python
- with some C code
- and some Cython code to tie everything together

# Cython you say?

```
from cpython.bytes cimport PyBytes_AS_STRING

cdef extern from "_hashindex.c":
    ctypedef struct HashIndex:
        pass

    void benchmark_getitem(HashIndex *index, char *keys, int key_count)

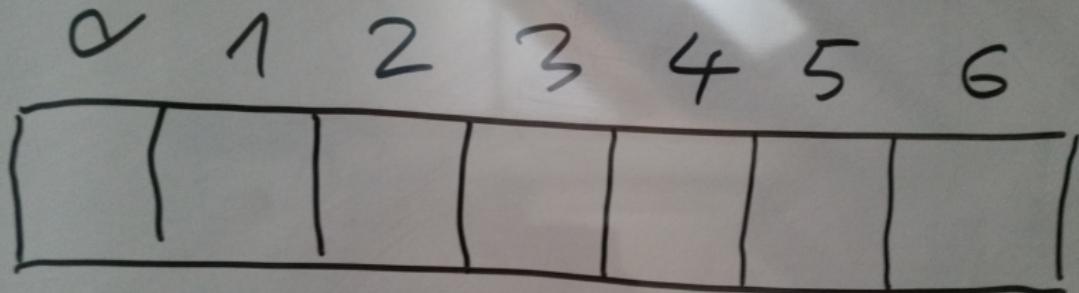
def bench_getitem(ChunkIndex chunk_index, bytes keys):
    cdef int key_count = len(keys) // chunk_index.key_size
    benchmark_getitem(
        chunk_index.index, PyBytes_AS_STRING(keys), key_count)
```

```
===== test session starts =====
platform linux -- Python 3.5.2, pytest-3.0.2, py-1.4.31, pluggy-0.3.1 -- /home/rctiorba/.venvs/borg/bin/python3
cachedir: .cache
benchmark: 3.0.0 (defaults: timer=time perf_counter disable_gc=False min_rounds=5 min_time=5.00us max_time=1.00s calibration_precision=10 warmup=False warmup_iterations=100000)
Tests enabled: BSD flags, nodes, atime/ntime, hardlinks, symlinks, root
Tests disabled: fuse
rootdir: /home/rctiorba/repos/borg, inifile: setup.cfg
plugins: cov-2.3.1, benchmark-3.0.0
collected 23 items

src/borg/testsuite/hashindex.py::HashIndexTestCase::test_chunkindex FAILED
src/borg/testsuite/hashindex.py::HashIndexTestCase::test_chunkindex_merge PASSED
src/borg/testsuite/hashindex.py::HashIndexTestCase::test_chunkindex_summarize PASSED
src/borg/testsuite/hashindex.py::HashIndexTestCase::test_iteritems PASSED
src/borg/testsuite/hashindex.py::HashIndexTestCase::test_nsindex FAILED
src/borg/testsuite/hashindex.py::HashIndexTestCase::test_resize fish "py test src/borg/testsuite/hash..." terminated by signal SIGSEGV (Address boundary error)
src/borg/testsuite/hashindex.py::HashIndexTestCase::test_resize_fish_terminated_by_signal_SIGSEGV_Address_boundary_error
```

# The goal: Robin Hood hashing

- Quick recap on hashmaps
- Open addressing
- Deletion



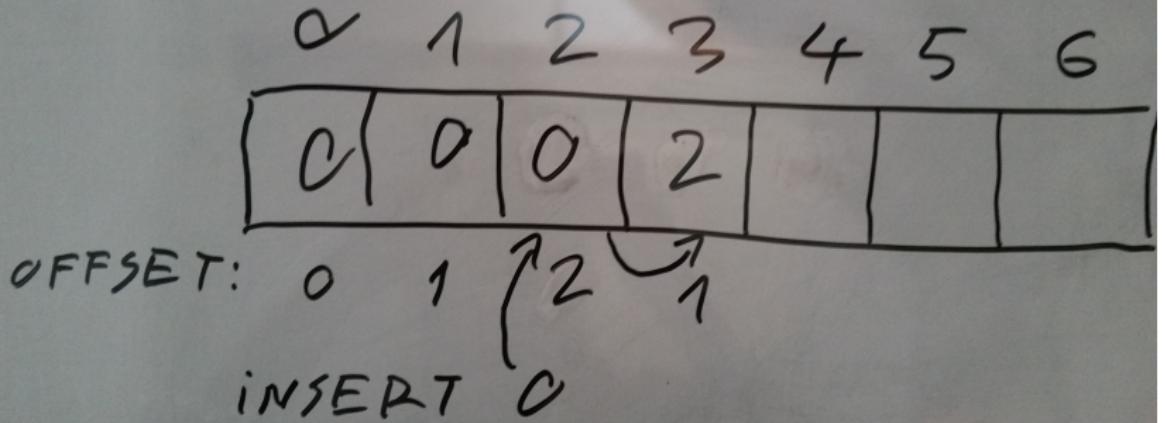
0	1	2	3	4	5	6
c	0	2	0	2		

	0	1	2	3	4	5	6
c	0	2					

0	1	2	3	4	5	6
c	0	2				

OFFSET: 0 1 0

INSERT 0



# Borg's current Hash

- key is 32 bytes, value is 12 bytes
- Open Addressing
- Tombstones
- hashindex\_lookup will promote buckets on top of tombstones

insert actual implementation here ...

# Cool story bro. Benchmarks?

```
def test_chunk_indexer_getitem(benchmark):
    max_key = 2**20
    index = ChunkIndex(max_key)
    keys = [sha256(H(k)).digest() for k in range(max_key)]
    for key in keys:
        index[key] = (0, 0, 0)

    def do_gets(keys=keys):
        for key in keys:
            index[key] # noqa

    benchmark.pedantic(do_gets, rounds=200)
```

# Cool story bro. Benchmarks?

## master with MAX\_HASH\_LOAD=0.93; 2^23 keys (reference)

Name (time in s)	Min	Max	Mean	benchmark: 2 tests			IQR
				StdDev	Median		
test_chunk_indexer_setitem	2.5804 (1.0)	2.6710 (1.0)	2.6278 (1.0)	0.0326 (6.19)	2.6251 (1.0)	0.0330	
test_chunk_indexer_getitem	4.4887 (1.74)	4.5013 (1.69)	4.4920 (1.71)	0.0053 (1.0)	4.4897 (1.71)	0.0039	

## robin\_hood with MAX\_HASH\_LOAD=0.93; 2^23 keys

Name (time in s)	Min	Max	Mean	benchmark: 2 tests			IQR
				StdDev	Median		
test_chunk_indexer_setitem	4.8533 (1.0)	4.8696 (1.0)	4.8591 (1.0)	0.0062 (1.0)	4.8569 (1.0)	0.0057	
test_chunk_indexer_getitem	4.9203 (1.01)	4.9791 (1.02)	4.9487 (1.02)	0.0211 (3.42)	4.9457 (1.02)	0.0218	

# What now?

- RH at least faster for missing items
- apply more coffee + time
- found a bug in hashindex\_set, wasn't ever doing the position swap

# Repeatable benchmarks

Name (time in ms)	Min	Max	Mean	benchmark: 2 tests	
				StdDev	Median
<hr/>					
test_chunk_indexer_setitem	112.1436 (1.0)	130.6243 (1.0)	116.1420 (1.0)	2.6196 (1.13)	115.1328 (1.0)
test_chunk_indexer_getitem	201.1404 (1.79)	216.5813 (1.66)	207.1470 (1.78)	2.3098 (1.0)	206.7306 (1.80)
<hr/>					
Name (time in ms)	Min	Max	Mean	benchmark: 2 tests	
				StdDev	Median
test_chunk_indexer_setitem	107.8818 (1.0)	122.5610 (1.0)	111.2042 (1.0)	2.3189 (1.0)	110.3162 (1.0)
test_chunk_indexer_getitem	197.3264 (1.83)	244.0130 (1.99)	203.2360 (1.83)	6.8099 (2.94)	202.0415 (1.83)

# Repeatable benchmarks

- Use same inputs
- Use same machine
- Make sure nothing else is running
- CPU frequency scaling?
- Isolate your code as much as possible

# Isolated your code for testing

```
def test_chunk_indexer_getitem(benchmark):
    max_key = 2**20
    index = ChunkIndex(max_key)
    keys = [sha256(H(k)).digest() for k in range(max_key)]
    for key in keys:
        index[key] = (0, 0, 0)

    def do_gets(keys=keys):
        for key in keys:
            index[key] # noqa

    benchmark.pedantic(do_gets, rounds=200)
```

# Isolated your code for testing

```
static void
benchmark_getitem(HashIndex *index, char *keys, int key_count)
{
    char *key = keys;
    char *last_addr = key + (32 * key_count);
    while (key < last_addr) {
        hashindex_get(index, key);
        key += 32;
    }
}
```

# Repeatable benchmarks

benchmark 'test_chunk_indexer_c_setitem_update': 6 tests						
Name (time in ms)	Min	Max	Mean	StdDev	Median	
test_chunk_indexer_c_setitem_update[0.3]	102.7014 (1.0)	102.9814 (1.0)	102.8013 (1.0)	0.0882 (1.0)	102.7865 (1.0)	
test_chunk_indexer_c_setitem_update[0.5]	173.7228 (1.69)	15,212.7702 (147.72)	2,462.7556 (23.96)	4,914.7807 (>1000.0)	235.4065 (2.29)	
test_chunk_indexer_c_setitem_update[0.86]	299.2544 (2.91)	27,044.4509 (262.61)	4,073.4831 (39.62)	8,542.5512 (>1000.0)	400.0225 (3.89)	
test_chunk_indexer_c_setitem_update[0.75]	324.0385 (3.16)	1,168.0475 (11.34)	549.7782 (5.35)	285.6380 (>1000.0)	420.7876 (4.09)	
test_chunk_indexer_c_setitem_update[0.93]	335.1741 (3.26)	4,270.4711 (41.47)	1,017.9627 (9.90)	1,459.4260 (>1000.0)	335.2918 (3.26)	
test_chunk_indexer_c_setitem_update[0.95]	344.4262 (3.35)	6,411.0975 (62.25)	951.2685 (9.25)	1,918.3883 (>1000.0)	344.6020 (3.35)	

benchmark 'test_chunk_indexer_c_setitem_update': 6 tests						
Name (time in ms)	Min	Max	Mean	StdDev	Median	
test_chunk_indexer_c_setitem_update[0.3]	102.7539 (1.0)	102.9468 (1.0)	102.8343 (1.0)	0.0602 (1.0)	102.8319 (1.0)	
test_chunk_indexer_c_setitem_update[0.5]	173.7920 (1.69)	15,279.5163 (148.42)	2,473.3088 (24.05)	4,937.4798 (>1000.0)	235.5947 (2.29)	
test_chunk_indexer_c_setitem_update[0.85]	299.0782 (2.91)	27,075.4213 (263.00)	4,077.7778 (39.65)	8,552.8954 (>1000.0)	400.2995 (3.89)	
test_chunk_indexer_c_setitem_update[0.75]	327.6863 (3.19)	1,167.8112 (11.34)	550.0308 (5.35)	285.1877 (>1000.0)	420.8317 (4.09)	
test_chunk_indexer_c_setitem_update[0.93]	334.9304 (3.26)	4,272.6914 (41.50)	1,018.8431 (9.91)	1,461.2820 (>1000.0)	335.2077 (3.26)	
test_chunk_indexer_c_setitem_update[0.95]	344.5489 (3.35)	6,424.4458 (62.41)	952.8873 (9.27)	1,922.5098 (>1000.0)	344.7731 (3.35)	

## We need a better way to present these benchmarks

- Robin Hood hashing first try
- Robin Hood, only check missing item shortcut every N buckets
- Robin Hood, shift entire chunk in one memmove
- Now let's replace modulo on master as well for apples to apples

# My takeaways

- MEASURE EVERYTHING!
- sometimes a 'worse' algorithm might be better
- integer division is slow

# Useful stuff

- Scott Meyers: Cpu Caches and Why You Care
- Andrei Alexandrescu: Writing Fast Code I
- pytest-benchmark
- kcachegrind + google-perf-tools + <https://pypi.python.org/pypi/yep>
- A raspberry pi for running the tests

```
> touch src/borg/hashindex.pyx ; pip install -e .. py test src/borg/testsuite/hashindex.py -v --capture=no -k test_chunkindex
Obtaining file:///home/rctorba/repos/borg
Requirement already satisfied: nsgpack-python>=0.4.6 in /home/rctorba/.venvs/borg/lib/python3.5/site-packages (from borgbackup==1.1.dev182+ng65ba75c.d20170216)
Installing collected packages: borgbackup
  Found existing installation: borgbackup 1.1.dev182+ng65ba75c d20170216
    Uninstalling borgbackup-1.1.dev182+ng65ba75c d20170216:
      Successfully uninstalled borgbackup-1.1.dev182+ng65ba75c d20170216
  Running setup.py develop for borgbackup
Successfully installed borgbackup
=====
===== test session starts =====
platform linux -- Python 3.5.2, pytest-3.0.2, py-1.4.31, pluggy-0.3.1 -- /home/rctorba/.venvs/borg/bin/python3
cachedir: .cache
benchmark: 3.0.0 (defaults: timer=time.perf_counter disable_gc=False min_rounds=5 min_time=5.00us max_time=1.00s calibration_precision=10 warmup=False warmup_iterations=1)
Testing BSD-style flags: yes
Testing fuse: no
fakeroot: no (>=1.20.2: no)
rootdir: /home/rctorba/repos/borg, infile: setup.cfg
plugins: cov-2.3.1, benchmark-3.0.0
collected 22 items

src/borg/testsuite/hashindex.py::HashIndexTestCase::test_chunkindex FAILED
src/borg/testsuite/hashindex.py::HashIndexTestCase::test_chunkindex_merge python3: malloc.c:2392: sysmalloc: Assertion `(old_top == initial_top (av) && old_size == 0) || (unsigned long) old_end & (pagesize - 1)) == 0)' failed.
  flush: "py test src/borg/testsuite/hash..." terminated by signal SIGABRT (Abort)
(borg)rctorba@void [-/borg] [key_bucket_split]
```

# Thanks

You can find the code here: <https://github.com/rciorba/borgbackup>

The slides are available at <https://devrandom.ro/talks>